

Lecture 2b: Informed Search

CSCI 360

Introduction to Artificial Intelligence

USC

What have we learned so far?

- What is AI?
- Problem-solving agent
- Uninformed search
- Informed search

What have we learned so far?

- **What is AI?**
 - Acting rationally
 - a function that maps percepts to actions
- Problem-solving agent
- Uninformed search
- Informed search

What have we learned so far?

- What is AI?
- **Problem-solving agent**
 - Formulating the problem
 - States, initial state, actions, transition model, goal test
- Uninformed search
- Informed search

What have we learned so far?


- What is AI?
- Problem-solving agent
- **Uninformed search**
 - Breadth-first search
 - Uniform-cost search
 - Depth-first search
 - Depth-limited search
 - Iterative deepening search
 - Bidirectional search
- Informed search

Recap: Breadth-first search (BFS)

- The “**shallowest**” node in frontier is chosen for expansion
 - The frontier is implemented using a FIFO queue


Recap: Breadth-first search (algorithm)

At each step, the node to be expanded next is indicated by the marker 




Recap: Breadth-first search (algorithm)

At each step, the node to be expanded next is indicated by the marker 




Recap: Breadth-first search (algorithm)

At each step, the node to be expanded next is indicated by the marker ►



Recap: Breadth-first search (algorithm)

At each step, the node to be expanded next is indicated by the marker ►




Recap: BFS (performance)

- Complete?
 - Yes -- If the shallowest goal node is at some finite depth, d , then BFS will eventually find the goal node
- Optimal?
 - Yes -- If the path cost is a non-decreasing function
- Time complexity?
 - $b + b^2 + b^3 + \dots + b^d = O(b^d)$
- Space complexity?
 - $O(b^d)$

Recap: Depth-first search (DFS)


- The “**deepest**” nodes in frontier is chosen for expansion
 - The frontier is implemented using a LIFO queue (or stack)



▶ -- node to be expanded next is indicated by the marker

Recap: Depth-first search (DFS)


- The “**deepest**” nodes in frontier is chosen for expansion
 - The frontier is implemented using a LIFO queue (or stack)



▶ -- node to be expanded next is indicated by the marker

Recap: Depth-first search (DFS)


- The “**deepest**” nodes in frontier is chosen for expansion
 - The frontier is implemented using a LIFO queue (or stack)



▶ -- node to be expanded next is indicated by the marker

Recap: Depth-first search (DFS)


- The “**deepest**” nodes in frontier is chosen for expansion
 - The frontier is implemented using a LIFO queue (or stack)



▶ -- node to be expanded next is indicated by the marker

Recap: Depth-first search (DFS)


- The “**deepest**” nodes in frontier is chosen for expansion
 - The frontier is implemented using a LIFO queue (or stack)



▶ -- node to be expanded next is indicated by the marker

Recap: Depth-first search (DFS)


- The “**deepest**” nodes in frontier is chosen for expansion
 - The frontier is implemented using a LIFO queue (or stack)



▶ -- node to be expanded next is indicated by the marker

Recap: Depth-first search (DFS)


- The “**deepest**” nodes in frontier is chosen for expansion
 - The frontier is implemented using a LIFO queue (or stack)



▶ -- node to be expanded next is indicated by the marker

Recap: Depth-first search (DFS)


- The “**deepest**” nodes in frontier is chosen for expansion
 - The frontier is implemented using a LIFO queue (or stack)



▶ -- node to be expanded next is indicated by the marker

Recap: Depth-first search (DFS)


- The “**deepest**” nodes in frontier is chosen for expansion
 - The frontier is implemented using a LIFO queue (or stack)



▶ -- node to be expanded next is indicated by the marker

Recap: Depth-first search (DFS)


- The “**deepest**” nodes in frontier is chosen for expansion
 - The frontier is implemented using a LIFO queue (or stack)



▶ -- node to be expanded next is indicated by the marker

Recap: Depth-first search (DFS)


- The “**deepest**” nodes in frontier is chosen for expansion
 - The frontier is implemented using a LIFO queue (or stack)



▶ -- node to be expanded next is indicated by the marker

Recap: Depth-first search (DFS)

- The “**deepest**” nodes in frontier is chosen for expansion
 - The frontier is implemented using a LIFO queue (or stack)



▶ -- node to be expanded next is indicated by the marker

Recap: DFS (performance)

- Complete?
 - No -- If the shallowest goal node is at some finite depth, ***d***, and the state space is infinite, then DFS may never find the goal node
- Optimal?
 - No
- Time complexity?
 - $b + b^2 + b^3 + \dots + b^m = O(b^m)$
- Space complexity?
 - $O(bm)$

Recap: DFS's space complexity

- Depth-first tree search is the “**basic workhorse**” of many areas of AI **because of its memory efficiency**
- Require storage of only **$O(bm)$** nodes

Depth	Nodes	Time	Memory	DFS
2	110	.11 milliseconds	107 kilobytes	
4	11,110	11 milliseconds	10.6 megabytes	
6	10^6	1.1 seconds	1 gigabyte	
8	10^8	2 minutes	103 gigabytes	
10	10^{10}	3 hours	10 terabytes	
12	10^{12}	13 days	1 petabyte	
14	10^{14}	3.5 years	99 petabytes	
16	10^{16}	350 years	10 exabytes	156 kilobytes

Figure 3.13 Time and memory requirements for breadth-first search. The numbers shown assume branching factor $b = 10$; 1 million nodes/second; 1000 bytes/node.

Recap: Having test of Both Worlds (DFS+BFS)?

- Iterative deepening DFS

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or failure
  for depth = 0 to  $\infty$  do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result
```



- Advantages
 - **Complete** and **Optimal** (just like BFS)
 - **O(bd)** storage requirement (just like DFS)

Recap: Iterative Deepening (example)

- Limit = 0

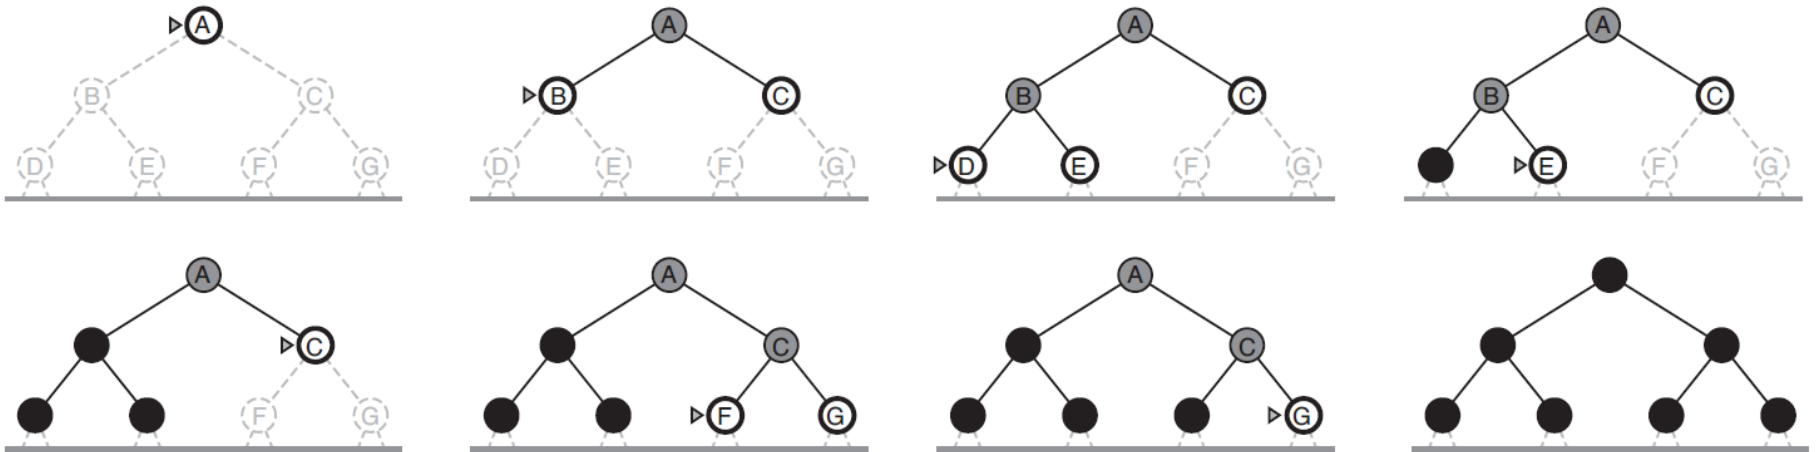


- Limit = 1




Recap: Iterative Deepening (example)

- Limit = 2



Recap: Iterative Deepening (example)




- Limit = 3



Recap: Iterative Deepening (Time Complexity)

- Except for (*limit* = *d*), each level is visited multiple times

$$N(\text{IDS}) = (d)b + (d - 1)b^2 + \dots + (1)b^d$$



Recap: Iterative Deepening (Time Complexity)

- Except for (*limit* = d), each level is visited multiple times

$$N(\text{IDS}) = (d)b + (d - 1)b^2 + \dots + (1)b^d$$

- Compare to BFS with shallowest goal depth (d) $\rightarrow O(b^d)$

- **Example:** Let ($b=10$) and ($d=5$)

$\begin{aligned} N(\text{IDS}) &= 50 + 400 + 3,000 + 20,000 + 100,000 = 123,450 \\ N(\text{BFS}) &= 10 + 100 + 1,000 + 10,000 + 100,000 = 111,110 \end{aligned}$
--

Recap: Iterative Deepening (conclusion)

- **Iterative deepening** is the preferred uninformed search method for many AI applications
 - E.g., when (1) the search space is large and (2) the depth of the solution is not known

Comparison of uninformed search algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ^a	Yes ^{a,b}	No	No	Yes ^a	Yes ^{a,d}
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$
Optimal?	Yes ^c	Yes	No	No	Yes ^c	Yes ^{c,d}

Figure 3.21 Evaluation of tree-search strategies. b is the branching factor; d is the depth of the shallowest solution; m is the maximum depth of the search tree; ℓ is the depth limit. Superscript caveats are as follows: ^a complete if b is finite; ^b complete if step costs $\geq \epsilon$ for positive ϵ ; ^c optimal if step costs are all identical; ^d if both directions use breadth-first search.

Outline for Today

- What is AI?
- Problem-solving agent
- Uninformed search
- **Informed search**
 - Greedy best-first search
 - A^* search
 - Memory-bounded heuristic search
 - Heuristic functions

Uninformed vs. Informed Search

- **Uninformed search** (or blind search): the agent does not know whether one state is “more promising”
 - All the agent can do is to (1) generate successor states and
 - (2) check if a state is a goal state
- **Informed Search**: the agent ***has some idea*** whether one state is “more promising” than another state

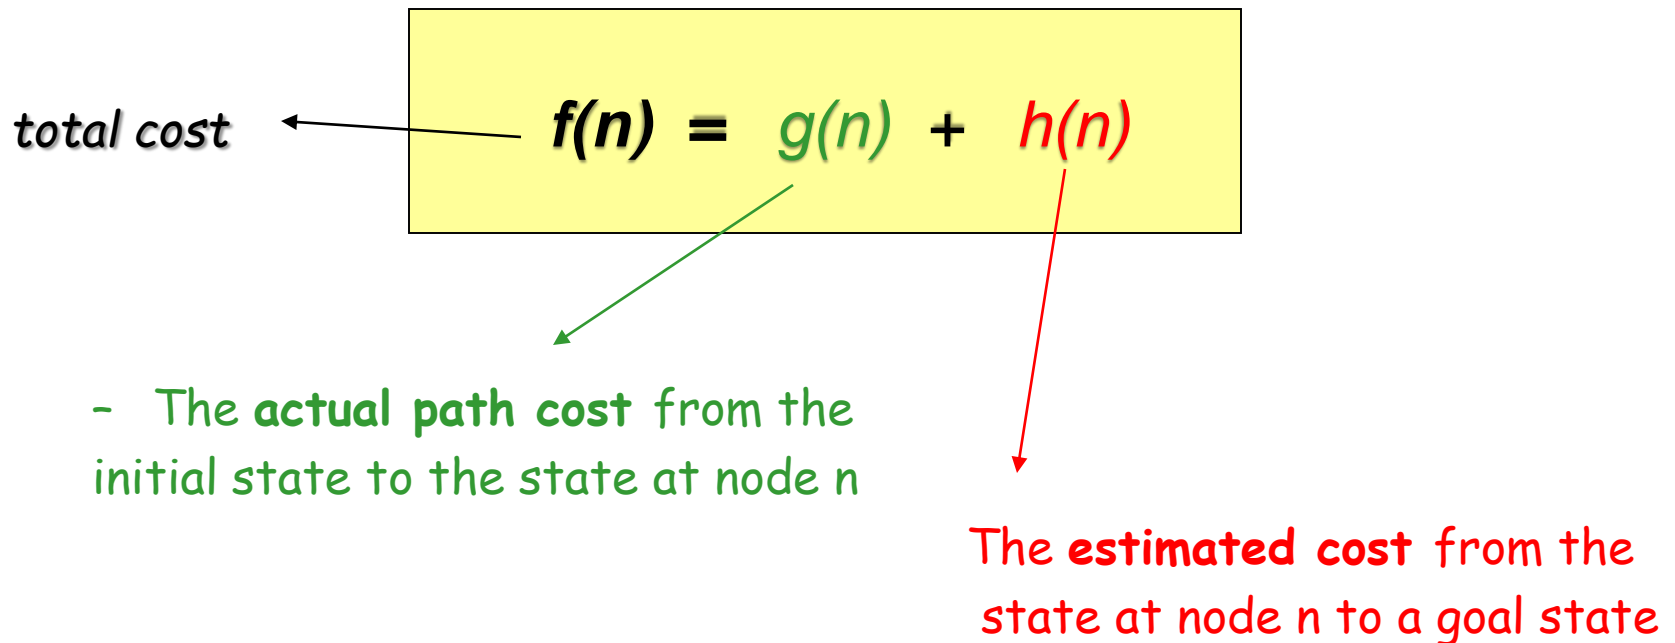
Best-First Search

- Node is chosen for expansion based on an **evaluation function, $f(n)$**
 - Returns an estimated, total solution cost at node n
- Pseudocode of “best-first search” is identical to “uniform-cost search” except that
 - **Uniform-cost**: priority queue ordered by actual path cost $g(n)$
 - **Best-first**: priority queue ordered by combined $f(n) = g(n) + h(n)$

What is $h(n)$?

Heuristic Function $h(n)$

- The estimated cost of the cheapest path from the state at node, n , to a goal state



Heuristic Function $h(n)$

- The estimated cost of the cheapest path from the state at node, n , to a goal state


$$f(n) = g(n) + h(n)$$

- Heuristic functions are not derived from the **problem description**, but from **“additional knowledge”**

Greedy Best-first Search

- **Intuition:** Expands node that “appears” to be **closest to the goal** since it’s likely to lead a solution quickly
 - i.e., $f(n) = h(n)$
 - In other words, assume the past cost $g(n) = 0$
- **Example:** In Map of Romania, if the goal is Bucharest, we can use the “*straight-line distance to Bucharest*” as the heuristic function $h(n)$

Question: *How to compute “straight-line distances”?*



Correct answer: You Can't!

Question: *How to compute the “straight-line distances”?*


The values of $h(n)$ cannot be computed from the problem description itself; it has to come from “**additional knowledge**”

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

Figure 3.22 Values of h_{SLD} —straight-line distances to Bucharest.

Greedy Best-first Search (example run)

- Starting at “Arad”, there are three actions
 - *Go(Zerind)*
 - *Go(Sibiu)*
 - *Go(Timisoara)*
- After that, which of the new nodes should be chosen for expansion?




Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

Greedy Best-first Search (example run)

- Starting at “Sibiu”, there are two actions

- *Go(Fagaras)*
- *Go(Rimnicu Vilcea)*




- After that, which of these two new nodes should be expanded next?

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374


Greedy Best-first Search (example run)

- Starting at “Fagaras”, there is only one action
 - *Go(Bucharest)*
- It's a goal state!
- But, is this the **optimal** solution?



Greedy Best-first Search (optimality)


- Starting at “Fagaras”, there is only one action
 - *Go(Bucharest)*
- It's a goal state!
- No, it is **not optimal**



278 kilometers vs. 310 kilometers

Greedy Best-first Search (completeness)

- Tree-Search is not complete even in a finite state space
- Example:
 - From “Iasi” to “Fagaras”
 - *Go(Neamt)*
 - *Go(Vaslui)*
 - Which node to expand next?
 - *Neamt*, since its “straight-line distance” to *Fagaras* is shorter
- Will never find the solution due to the infinite loop “Neamt - Iasi”



Graph-Search is complete, but only in a finite state space


Greedy Best-first Search (example run)

- When it works, it's very fast due to “pruning” of state space



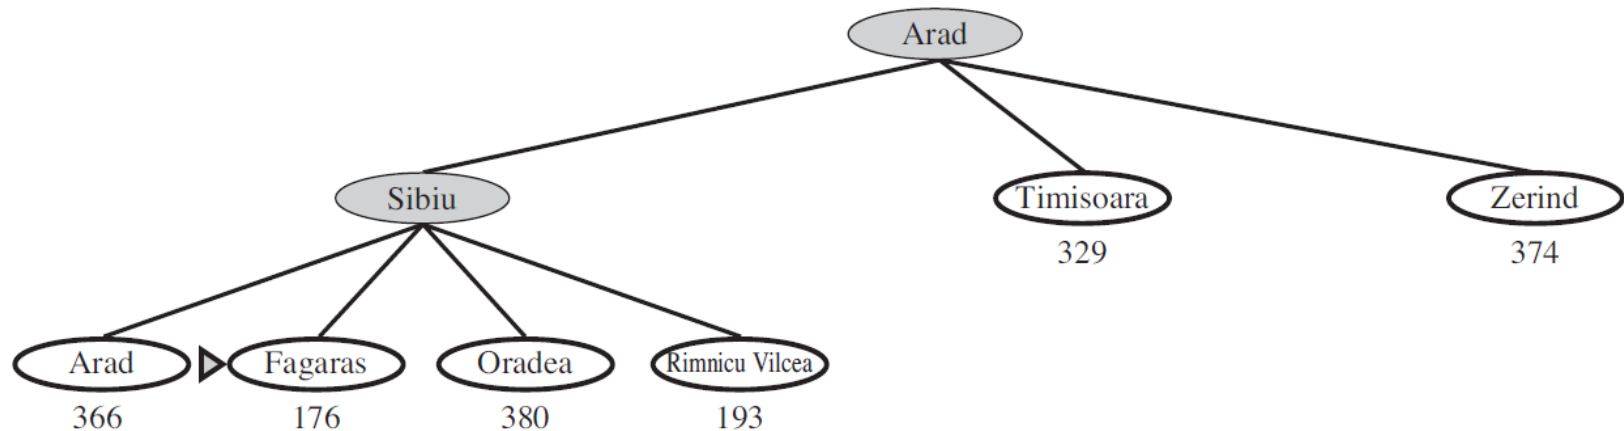
Greedy Best-first Search (example run)

- When it works, it's very fast due to “pruning” of state space



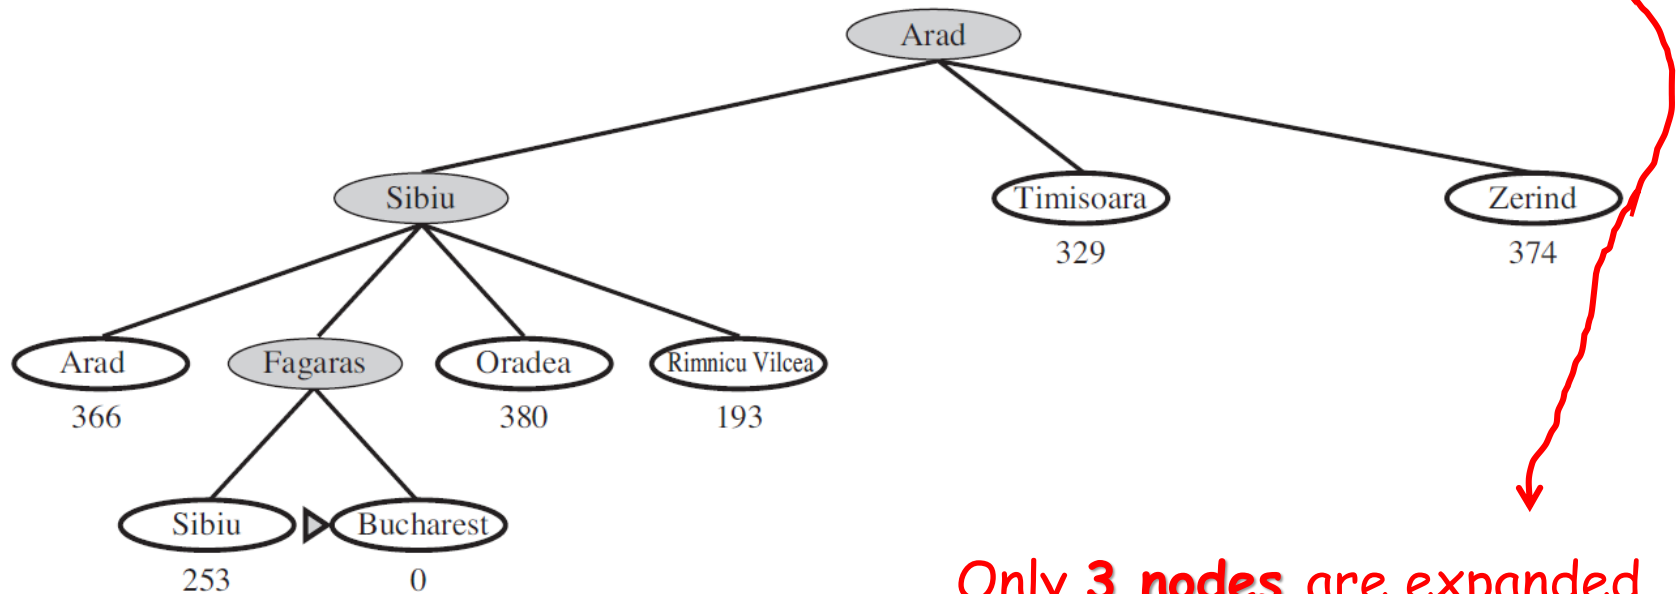
Greedy Best-first Search (example run)

- When it works, it's very fast due to “pruning” of state space



Greedy Best-first Search (example run)

- When it works, it's very fast due to 'pruning' of state space



Only **3 nodes** are expanded
(Arad, Sibiu, Fagaras)

Timisoara is never expanded

Summary of Greedy Best-first Search

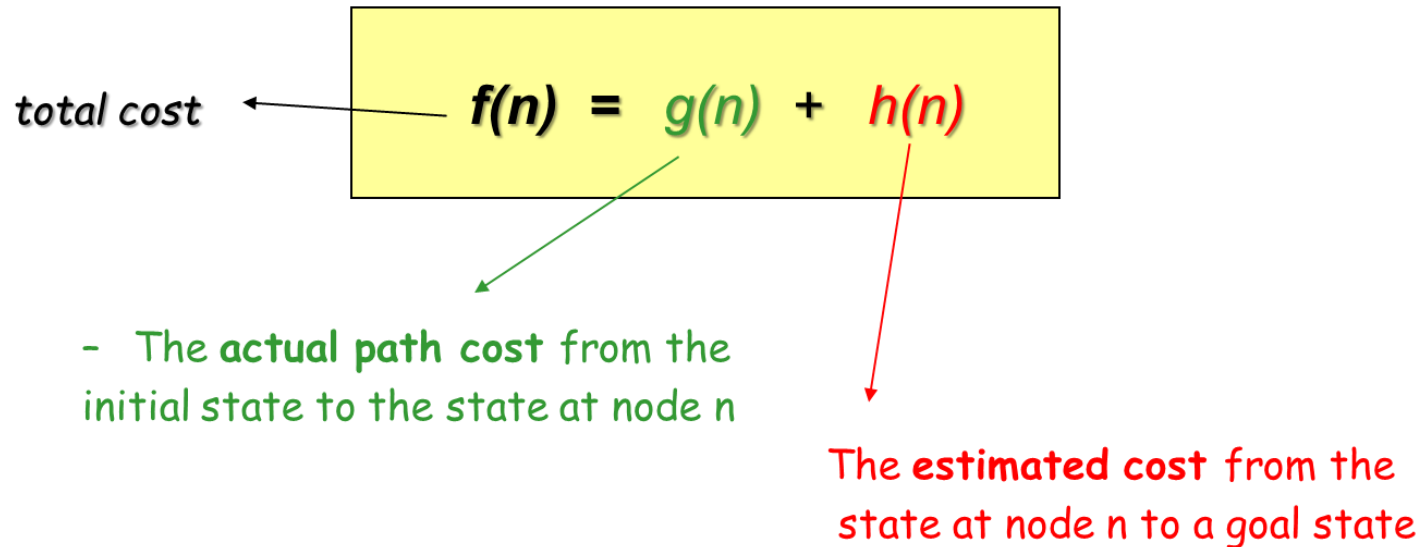
- Complete?
 - No – can get stuck in loops
- Optimal?
 - No
- Time?
 - Exponential, but a way to get dramatic improvement
- Space?
 - Keeps all nodes in memory

Outline for Today

- What is AI?
- Problem-solving agent
- Uninformed search
- **Informed search**
 - Greedy best-first search
 - **A* search**
 - Memory-bounded heuristic search
 - Heuristic functions

Question: How to make “best-first search” optimal and complete?

- **Answer:** Use both $g(n)$ and $h(n)$ to compute total cost $f(n)$




A* search — minimizing the total estimated cost

- Most widely known form of “best-first search”
 - **Goal:** find the cheapest solution
 - **At each step:** expand node with lowest value of the total estimated solution cost: $f(n) = g(n) + h(n)$
- Pseudo code is identical to “Uniform-cost Search” except for using **$(g+h)$** instead of **(g)** to set the priority queue

Key idea: Avoid expanding along paths that are already known to be expensive


A* search (example run)

- Total estimated solution cost: $f(n) = g(n) + h(n)$



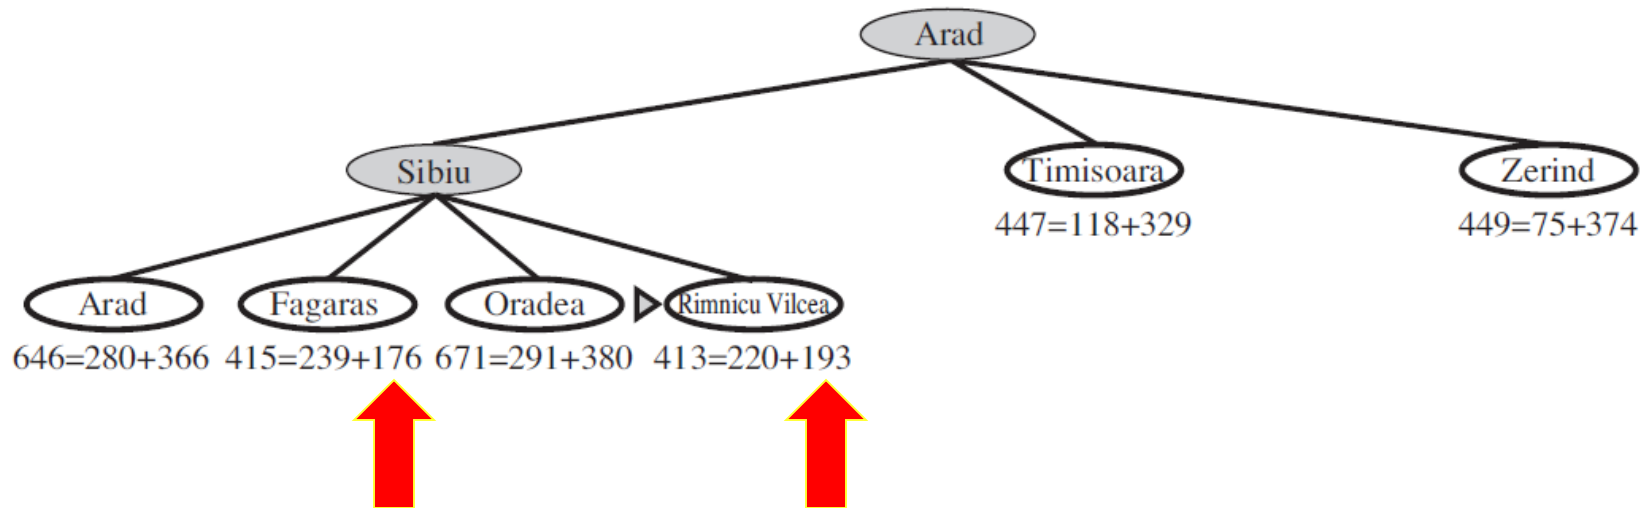
A* search (example run)

- Total estimated solution cost: $f(n) = g(n) + h(n)$



A* search (example run)

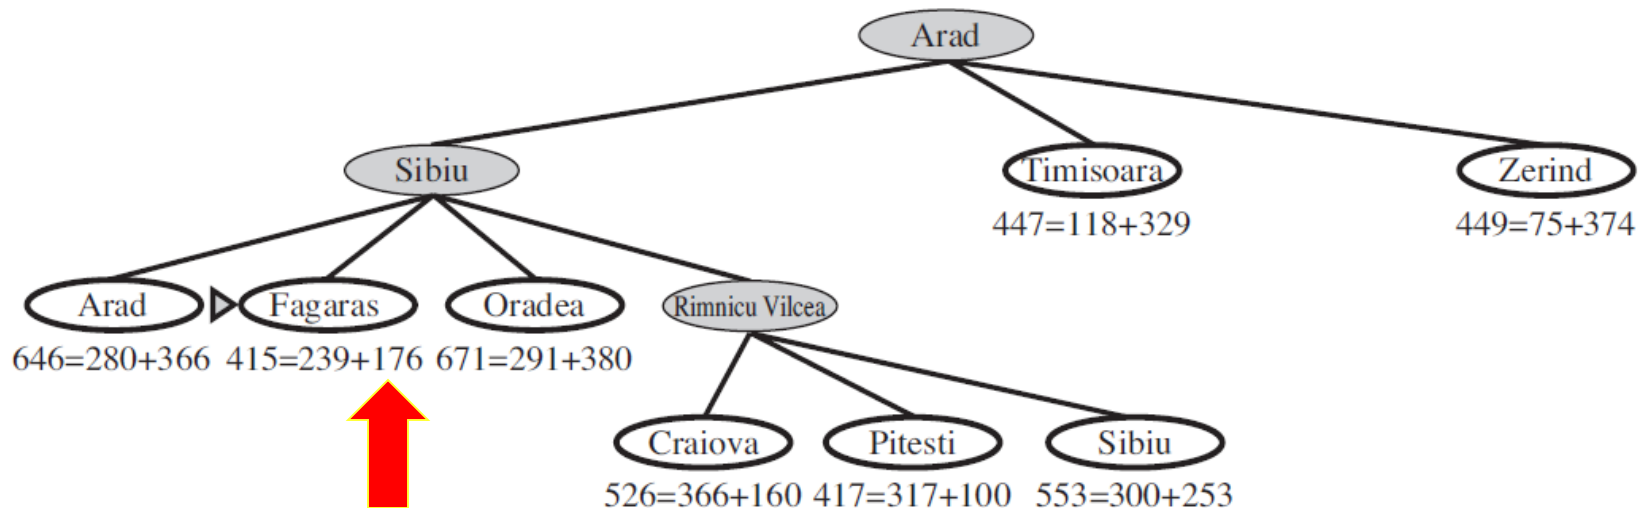
- Total estimated solution cost: $f(n) = g(n) + h(n)$



Difference between A* search
and Greedy best-first search

A* search (example run)

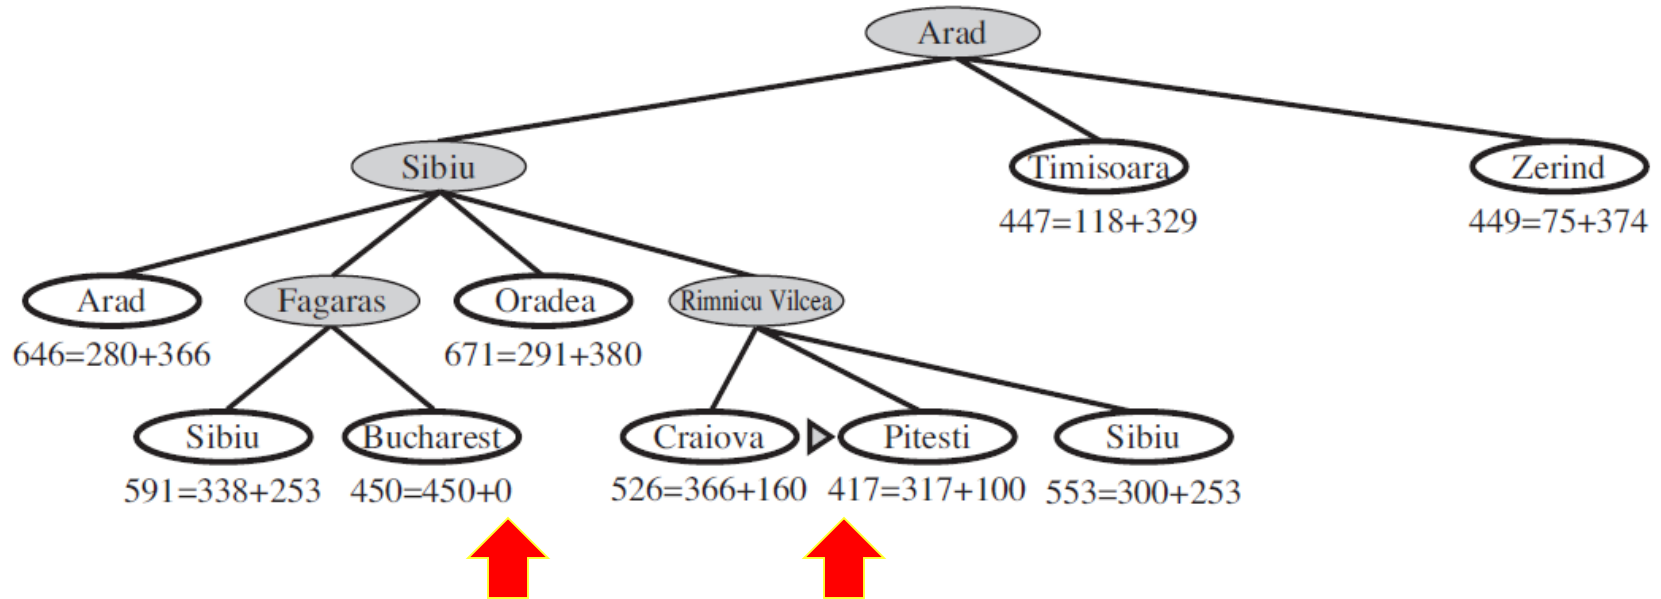
- Total estimated solution cost: $f(n) = g(n) + h(n)$



Difference between A* search
and Greedy best-first search

A* search (example run)

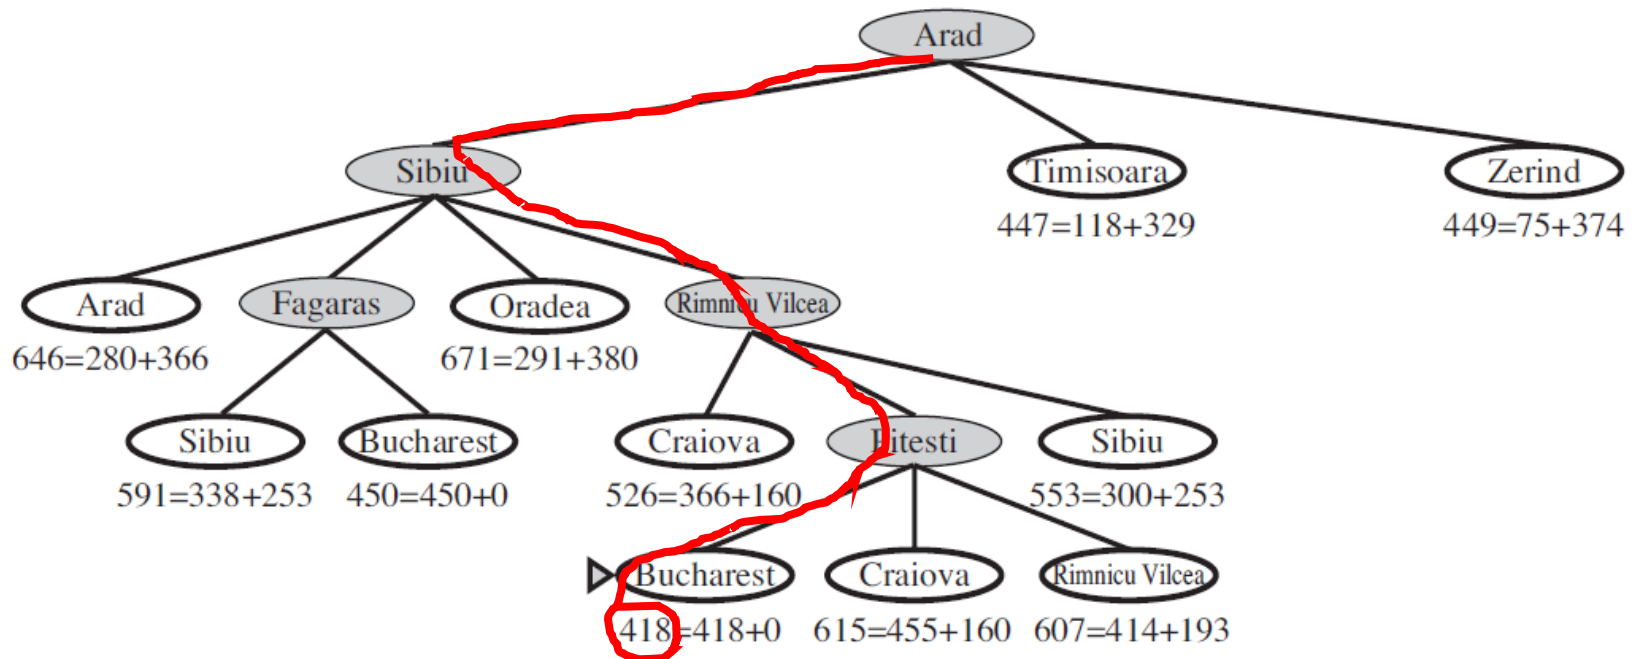
- Total estimated solution cost: $f(n) = g(n) + h(n)$



Greedy best-first search would have stopped at this point

A* search (example run)

- Total estimated solution cost: $f(n) = g(n) + h(n)$



Optimal solution

Pseudo code (same as Uniform-cost Search)

- Node with the “**lowest path cost**” → node with the “**lowest total cost**”

```
function UNIFORM-COST-SEARCH(problem) returns a solution, or failure
  node ← a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  frontier ← a priority queue ordered by PATH-COST, with node as the only element
  explored ← an empty set
  loop do
    if EMPTY?(frontier) then return failure
    node ← POP(frontier) /* chooses the lowest-cost node in frontier */
    if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child ← CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        frontier ← INSERT(child, frontier)
      else if child.STATE is in frontier with higher PATH-COST then
        replace that frontier node with child
```

Is A^* complete and optimal?


- Complete
 - If there is a solution, A^* search will find it
- Optimal
 - The solution found by A^* search is guaranteed to be the cheapest solution

Condition for optimality

- Heuristic function $h(n)$ is **admissible**
- An admissible heuristic is one that *never overestimates* the true cost to reach the goal
 - Since $f(n) = g(n) + h(n)$ and
 - $g(n)$ is the actual cost,
 - If $h(n)$ never overestimates the cost from n to the goal, then $f(n)$ never overestimates the total cost of a solution along the current path through node n either

Stronger condition for optimality

- **Consistency** (or **monotonicity**) is required only for applications of A^* to **graph search**
- Heuristic function $h(n)$ is consistent if, for any successor n' of n generated by any action a , we have
 - $h(n) \leq c(n,a,n') + h(n')$



Optimality of A^*


- The graph-search version of A^* search is optimal if $h(n)$ is consistent.
- (Less useful) proof sketch:
 - First, if $h(n)$ is consistent, the values of $f(n)$ along any path are non-decreasing.
 - By the definition of “consistency”
 - Second, whenever A^* selects a node (n) for expansion, the optimal path to that node has been found
 - Assume this is not true, there would have to be another frontier node (n') on the optimal path from the start node to n ; since $f(n)$ is non-decreasing, $f(n')$ would have a lower value than $f(n)$, which means n' would have been selected for expansion

Reminder: *What's the meaning of $f(n)$?*

For some admissible heuristics, f may *decrease* along a path

E.g., suppose n' is a successor of n

(not consistent)




Reminder: *What's the meaning of $f(n)$?*

For some admissible heuristics, f may *decrease* along a path

E.g., suppose n' is a successor of n

(not consistent)



But this throws away information!

$f(n) = 9 \Rightarrow$ true cost of a path through n is ≥ 9


Hence true cost of a path through n' is ≥ 9 also

Reminder: *What's the meaning of $f(n)$?*

For some admissible heuristics, f may *decrease* along a path

E.g., suppose n' is a successor of n

(not consistent)



But this throws away information!

$f(n) = 9 \Rightarrow$ true cost of a path through n is ≥ 9

Hence true cost of a path through n' is ≥ 9 also


Pathmax modification to A^* :

Instead of $f(n') = g(n') + h(n')$, use $f(n') = \max(g(n') + h(n'), f(n))$

With pathmax, f is always nondecreasing along any path

Contours in A* search


- The values of $f(n)$ are non-decreasing along any path from the initial state – they form “contours”



Question: Are there contours in “uniform-cost search”? What do they look like?

Contours in Uniform-cost Search

- They are “circular bands” around the start state

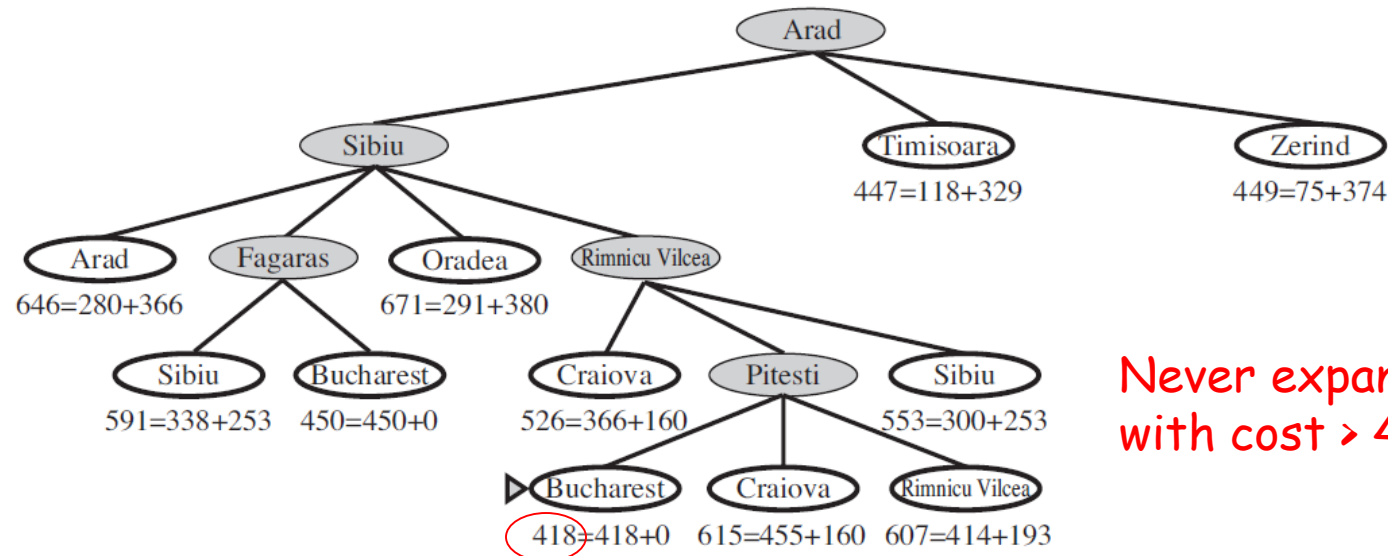


Properties of A^* (summary)

- Complete?
 - Yes, unless infinitely many nodes with $f \leq f(G)$
- Optimal?
 - Yes – it cannot expand contour f_{i+1} until f_i is finished
- Time?
 - Exponential
- Space?
 - Keep all nodes in memory

A* is also “optimally efficient”

- No other optimal algorithm is guaranteed to expand fewer nodes than A*
 - First, A* expands no nodes with $f(n) > C^*$, where C^* is the cost of the optimal solution
 - Second, any other algorithm that does not expand all nodes with $f(n) < C^*$ runs the risk of missing the optimal solution



Never expand nodes
with cost > 418

Space complexity of A^* is a bigger problem

- A^* usually runs out of space long before it runs out of time
 - Because it keeps all generated nodes in memory
 - And the number of generated nodes is exponential in C^*
- Solution?
 - Memory-bounded heuristic search

Outline for Today

- What is AI?
- Problem-solving agent
- Uninformed search
- **Informed search**
 - Greedy best-first search
 - A^* search
 - **Memory-bound heuristic search**
 - Heuristic functions

Three variants

- Iterative-deepening A^* (IDA*)
- Recursive best-first search (RBFS)
- Simplified memory-bound A^* (SMA*)

Three variants

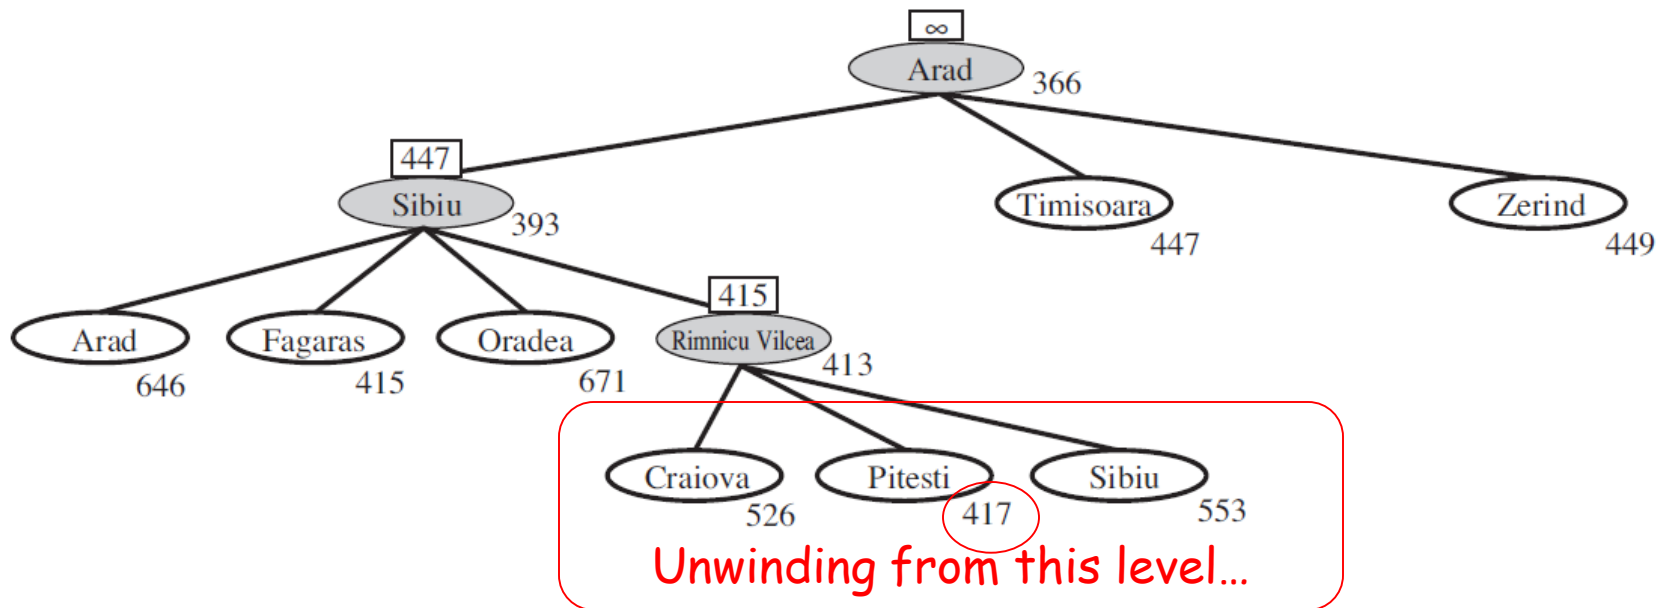
- Iterative-deepening A^* (IDA*)
 - Instead of using “depth” as the cutoff, IDA* uses the f-cost ($g+h$) as the cutoff
 - No need to store the sorted queue of nodes
- Recursive best-first search (RBFS)
- Simplified memory-bound A^* (SMA*)

Three variants

- Iterative-deepening A^* (IDA*)
- Recursive best-first search (RBFS)
 - Similar to recursive DFS with depth limit, but uses *f-limit* to keep track of the total cost of the best alternative path available from any ancestor of the current node
 - If the cost of the current node (*f-cost*) exceeds this limit (*f-limit*), the recursion unwinds back to the alternative path
- Simplified memory-bound A^* (SMA*)

RBFS (recursive best-first search)

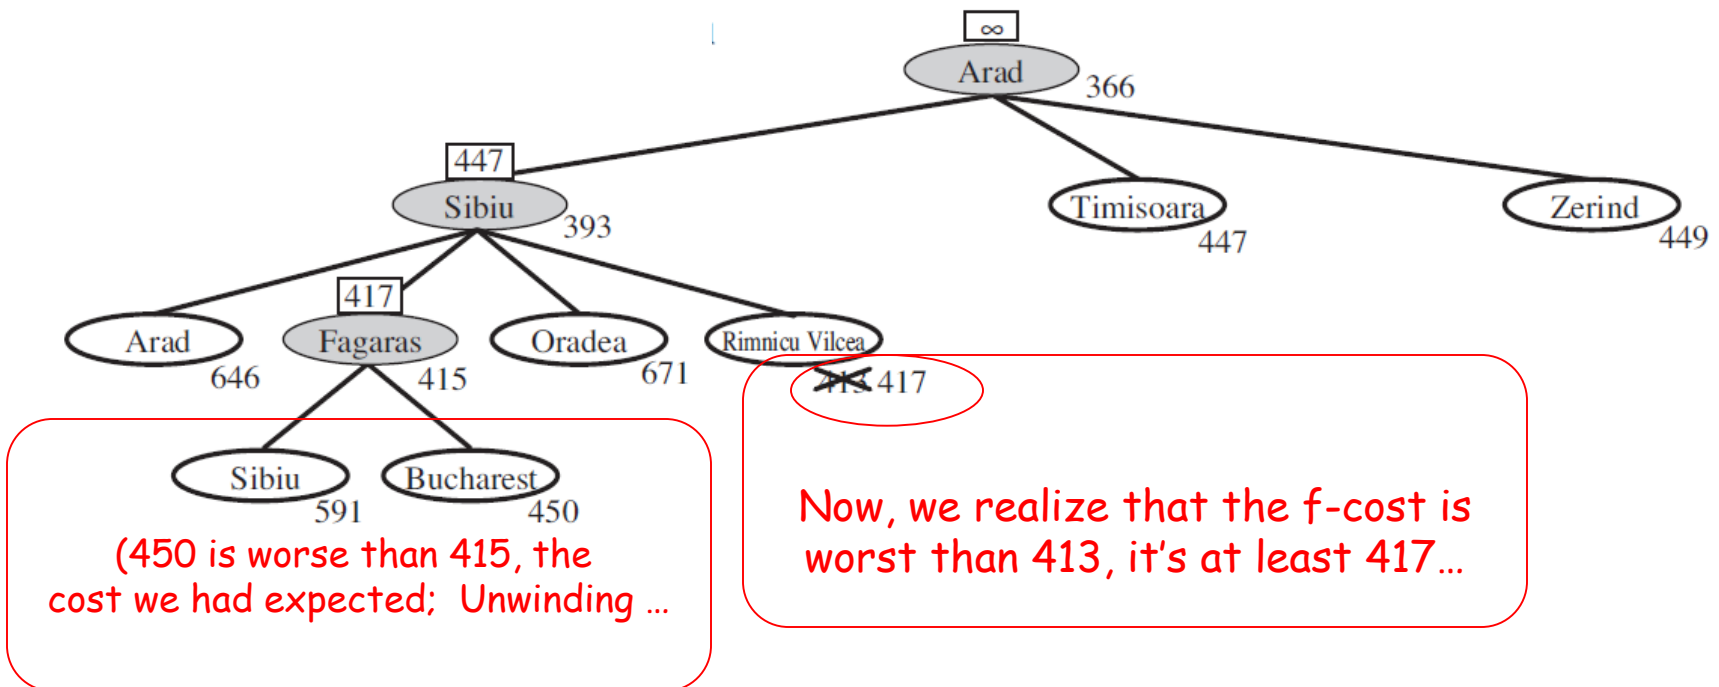
- **Above a node:** the f-limit for each recursive call on the node
- **Under a node:** f-cost associated with the node



- Before expanding "Sibiu", it computes the cheapest alternative path from sibling cities "Timisoara" and "Zerind" (and the f-limit is 447)
- Before expanding "Rimnicu Vilcea", it compute the cheapest alternative path from sibling cities "Arad", "Fargaras" and "Oradea" (and the f-limit is 415)

RBFS (recursive best-first search)

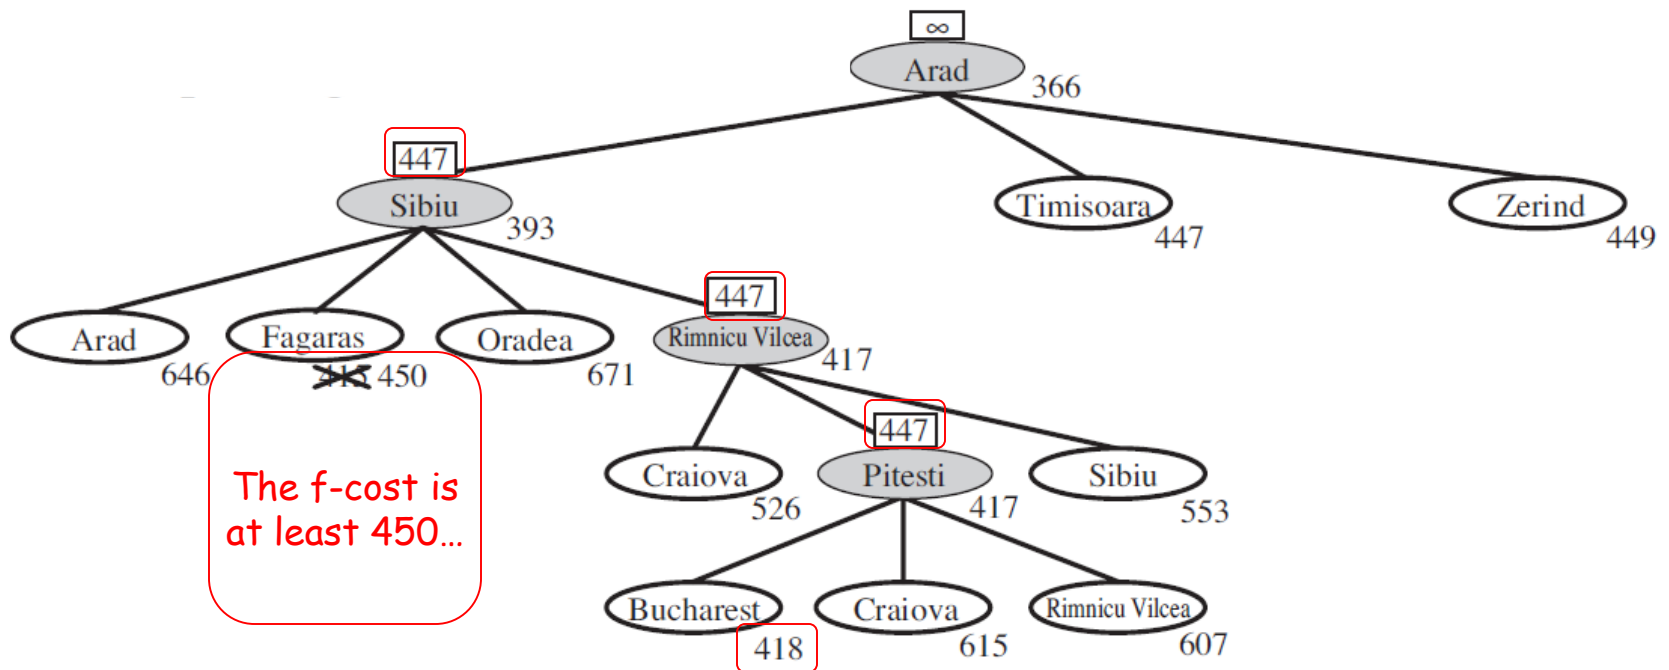
- **Above a node:** the f-limit for each recursive call on the node
- **Under a node:** f-cost associated with the node



- Before expanding "Fagaras", the f-limit is set to 417 (newly computed from "Rimnicu Vilcea")
- When "Bucharest" has f-cost=450, which is greater than f-limit=417, it unwinds again...

RBFS (recursive best-first search)

Question: why is this "space efficient"?



Performance of RBFS

- Space complexity is **linear** in the depth of the deepest optimal solution
- **Optimal** if the heuristic function $h(n)$ is admissible
- Time complexity is rather difficult to characterize
 - Depending on the accuracy of the heuristic function and how often the best path changes as the nodes are expanded

Three variants

- Iterative-deepening A^* (IDA*)
- Recursive best-first search (RBFS)
- Simplified memory-bound A^* (SMA*)
 - SMA* starts just like A^* , expanding the best leaf until memory is full
 - At this point, it adds a new node to the search tree only after dropping an old one (the worst leaf node, with the highest f-cost)

Performance of SMA*

- Complete if there is any reachable solution
 - That is, the depth of the shallowest goal node is less than the memory size
- Optimal if any optimal solution is reachable
 - Otherwise, it returns the best reachable solution

Outline for Today

- What is AI?
- Problem-solving agent
- Uninformed search
- **Informed search**
 - Greedy best-first search
 - A^* search
 - Memory-bound heuristic search
 - **Heuristic functions**

How to evaluate heuristic functions?

- Example: 8-puzzle
 - Average solution cost is 22 steps
 - Branching factor < 3
 - Exhaustive tree search to depth 22 $\rightarrow 3^{22}$ states
 - Graph search $\rightarrow 181,440$ distinct states

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Heuristic functions for 8-puzzle

- Two candidates
 - $h1(n)$ = number of misplaced tiles
 - $h2(n)$ = sum of the distances of the tiles from their goal positions
 - Manhattan distance: The sum of the horizontal and vertical distances

$h1 = ??$

$h2 = ??$

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Heuristic functions for 8-puzzle

- Two candidates
 - $h1(n)$ = number of misplaced tiles
 - $h2(n)$ = sum of the distances of the tiles from their goal positions
 - Manhattan distance: The sum of the horizontal and vertical distances

$h1 = 8$

$h2 = 18$

How do you
know which
heuristics is
better?

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Experimental results: A* with h_1 and h_2

- Solving 1200 random 8-puzzle problems

	Search Cost (nodes generated)		
d	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6
4	112	13	12
6	680	20	18
8	6384	39	25
10	47127	93	39
12	3644035	227	73
14	—	539	113
16	—	1301	211
18	—	3056	363
20	—	7276	676
22	—	18094	1219
24	—	39135	1641

The effect of heuristic accuracy

- Effective branching factor b^*
 - Assume the heuristic is perfect, A^* would explore a total of (N) nodes and the solution depth is d

$$N+1 = 1 + b^* + (b^*)^2 + \dots + (b^*)^d$$

Let $(N=52)$ and $(d=5)$, what would be the value of (b^*) ?

Answer: $b^*=1.92$

Experimental results: A* with h1 and h2

- Solving 1200 random 8-puzzle problems

	Search Cost (nodes generated)			Effective Branching Factor		
d	IDS	$A^*(h_1)$	$A^*(h_2)$	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	3644035	227	73	2.78	1.42	1.24
14	–	539	113	–	1.44	1.23
16	–	1301	211	–	1.45	1.25
18	–	3056	363	–	1.46	1.26
20	–	7276	676	–	1.47	1.27
22	–	18094	1219	–	1.48	1.28
24	–	39135	1641	–	1.48	1.26

Why Heuristic $h_2(n)$ is better than $h_1(n)$?

- By their definitions, for any node n , we have $h_1(n) \leq h_2(n)$
 - We say that h_2 dominates h_1
- **Domination translates into efficiency:** A^* using h_2 will never expand more nodes than A^* using h_1

$$f(n) < C^*$$

$$g(n) + h(n) < C^*$$

$$g(n) < C^* - h(n)$$

Why Heuristic $h_2(n)$ is better than $h_1(n)$?

- By their definitions, for any node n , we have $h_1(n) \leq h_2(n)$
 - We say that h_2 dominates h_1
- **Domination translates into efficiency:** A^* using h_2 will never expand more nodes than A^* using h_1

$$f(n) < C^*$$

$$g(n) + h(n) < C^*$$

$$g(n) < C^* - h(n)$$

larger $h(n) \rightarrow$ smaller $g(n)$

when $h(n)=0$, $g(n)$ is only bounded by C^* (i.e., no guidance from the goal at all)

How to generate heuristic functions?

- Admissible heuristics can be derived from the exact solution cost of a “**relaxed**” version of the problem
- By “**relaxed**” we mean it has an “**over-approximation**” of the transition model of the original problem
 - More edges between states, to provide short cuts

Relaxed problems for 8-puzzle

- Original transition model

A tile can move from square A to square B if
A is horizontally or vertically adjacent to B, and B is blank

- Relaxed transition model

- (a) A tile can move from square A to square B if A is adjacent to B
- (b) A tile can move from square A to square B if B is blank
- (c) A tile can move from square A to square B.

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Relaxed problems for 8-puzzle

- Original transition model

A tile can move from square A to square B if
A is horizontally or vertically adjacent to B, and B is blank

- Relaxed transition model

- (a) A tile can move from square A to square B if A is adjacent to B
- (b) A tile can move from square A to square B if B is blank
- (c) A tile can move from square A to square B.

Number of misplaced tiles

Manhattan distance

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Combining heuristic functions

- Given a collection of admissible heuristics, the composite heuristic (which uses whichever function that is the most accurate on the node in question) is also admissible

$$h(n) = \max\{h_1(n), \dots, h_m(n)\}$$

- Furthermore, $h(n)$ dominates all of these component heuristics

Generating admissible heuristics from subproblems

- **Pattern Database:** store the exact solution costs for every possible subproblem instance
 - Every possible configuration of the first four tiles and the blank
 - Locations of the other four tiles are irrelevant (but moving them still counts toward the solution cost)
- The database can be constructed by searching back from the goal, and recording the cost of each pattern encountered

*	2	4
*		*
*	3	1

Start State

	1	2
3	4	*
*	*	*

Goal State

Learning heuristic functions from experience

- Inductive learning of a heuristic function **$h(n)$** in terms of features **$x_1(n)$** and **$x_2(n)$**

$$h(n) = c_1x_1(n) + c_2x_2(n)$$

- Candidate features
 - “number of misplaced tiles”
 - “number of pairs of adjacent tiles that are not adjacent in the goal state”

Outline for Today

- What is AI?
- Problem-solving agent
- Uninformed search
- **Informed search**
 - Greedy best-first search
 - A^* search
 - Memory-bound heuristic search
 - Heuristic functions
 - **Summary of today's lecture**

Summary

- Informed search may have access to a heuristic function, $h(n)$, which estimate the cost of a solution from n
 - Greedy best-first search expands nodes with minimal $h(n)$. It is not optimal but is often fast
 - A^* search expands nodes with minimal $f(n)=g(n)+h(n)$. It is complete and optimal, provided that $h(n)$ is admissible (or consistent). The space complexity of A^* is still high.
 - RBFS and SMA* are robust, optimal algorithms that use limited memory
- The performance of heuristic search algorithms depends on the quality of the heuristic function
 - Good heuristics can be constructed by (1) relaxing the problem definition, (2) storing precomputed solution costs for subproblems in a pattern database, or by (3) learning from experience