

# Lecture 15b: Reinforcement Learning

---


CSCI 360

Introduction to Artificial Intelligence

USC

# Here is where we are...


|    |              |              |  |   |
|----|--------------|--------------|--|---|
|    | 3/1          |              | Project 2 Out  |   |
| 9  | 3/4<br>3/6   | 3/5<br>3/7   | Quantifying Uncertainty<br>Bayesian Networks   | [Ch 13.1-13.6]<br>[Ch 14.1-14.2]          |
| 10 | 3/11<br>3/13 | 3/12<br>3/14 | (spring break, no class)<br>(spring break, no class)   |   |
| 11 | 3/18<br>3/20 | 3/19<br>3/21 | Inference in Bayesian Networks<br>Decision Theory  | [Ch 14.3-14.4]<br>[Ch 16.1-16.3 and 16.5] |
|    | 3/23         |              | Project 2 Due  |   |
| 12 | 3/25<br>3/27 | 3/26<br>3/28 | <i>Advanced topics (Chao traveling to National Science Foundation)</i><br><i>Advanced topics (Chao traveling to National Science Foundation)</i> |   |
|    | 3/29         |              | Homework 2 Out   |   |
| 13 | 4/1<br>4/3   | 4/2<br>4/4   | Markov Decision Processes<br>Decision Tree Learning  | [Ch 17.1-17.2]<br>[Ch 18.1-18.3]          |
|    | 4/5<br>4/5   |              | Homework 2 Due<br>Project 3 Out  |   |
| 14 | 4/8<br>4/10  | 4/9<br>4/11  | Perceptron Learning<br>Neural Network Learning   | [Ch 18.6]<br>[Ch 18.7]                    |
| 15 | 4/15<br>4/17 | 4/16<br>4/18 | Statistical Learning<br>Reinforcement Learning   | [Ch 20.2.1-20.2.2]<br>[Ch 21.1-21.2]      |
| 16 | 4/22<br>4/24 | 4/23<br>4/25 | Artificial Intelligence Ethics<br>Wrap-Up and Final Review   |   |
|    | 4/26         |              | Project 3 Due  |   |
|    | 5/3          | 5/2          | Final Exam (2pm-4pm)   |   |



# Outline

---


- What is AI?
- Part I: Search
- Part II: Logical reasoning
- Part III: Probabilistic reasoning
- **Part IV: Machine learning**
  - Decision Tree Learning
  - Perceptron Learning
  - Neural Network Learning
  - Statistical Learning
  - **Reinforcement Learning**



# Example **search** problem



---

- Initial state: (1,1)
- Goal state: (4,3) utility +1  
(4,2) utility -1
- Actions:
  - Up, Down, Left, Right reward -0.04
  - *(won't move it running into the wall)*
- Transition model:
  - If  $\text{RESULT}(s, a)$  is **deterministic**, then it's a **search** problem





# Example adversarial search problem

- Initial state: (1,1)
- Goal state: (4,3) utility +1  
(4,2) utility -1
- Actions:
  - Up, Down, Left, Right reward -0.04
  - (won't move it running into the wall)
- Transition model:
  - If  $\text{RESULT}(s, a)$  is non-deterministic, then it's an adversarial search problem





# Example **MDP** (Markov Decision Process) problem

- Initial state:  $(1,1)$
- Goal state:  $(4,3)$  utility  $+1$   
 $(4,2)$  utility  $-1$
- Actions:
  - Up, Down, Left, Right reward  $-0.04$
  - (won't move it running into the wall)*
- Transition model:
  - If  $\text{RESULT}(s, a)$  is **non-deterministic**,  
and **probabilistic**, then it's **MDP**  
(Markov decision process)



# Example reinforcement learning problem

- Initial state: (1,1)
- Goal state: (4,3) utility +1  
(4,2) utility -1
- Actions:
  - Up, Down, Left, Right reward -0.04
  - ~~(won't move it running into the wall)~~
- Transition model:
  - If  $\text{RESULT}(s, a)$  is unknown, then it's reinforcement learning



# Markov decision process (MDP)

---

- MDP is a sequential decision problem that has
  - a **fully observable, stochastic** environment
  - a **Markovian** transition model, and
  - the **additive** rewards
- State space formalism
  - Initial state ( $s_0$ )
  - $\text{Actions}(s) = \{a_1, a_2, \dots\}$  for each state
  - Transition model  $P(s' \mid s, a)$  for each state and each action
  - Reward function  $R(s)$



# Reinforcement learning: *Compared to MDP*

---

- MDP is a sequential decision problem that has
  - a **fully observable, stochastic** environment
  - ~~– a **Markovian** transition model, and~~
  - the **additive** rewards
- State space formalism
  - Initial state ( $s_0$ )
  - Actions( $s$ ) =  $\{a_1, a_2, \dots\}$  for each state
  - ~~– Transition model  $P(s' | s, a)$  for each state and each action~~
  - ~~– Reward function  $R(s)$~~

# The Story So Far: MDPs and RL

## Known MDP: Offline Solution

### Goal

Compute  $U^*, Q^*, \pi^*$

Evaluate a fixed policy  $\pi$

### Technique

Value / policy iteration

Policy evaluation

## Unknown MDP: Model-Based

### Goal

Compute  $U^*, Q^*, \pi^*$

Evaluate a fixed policy  $\pi$

### Technique

VI/PI on approx. MDP

PE on approx. MDP

## Unknown MDP: Model-Free

### Goal

Compute  $U^*, Q^*, \pi^*$

Evaluate a fixed policy  $\pi$

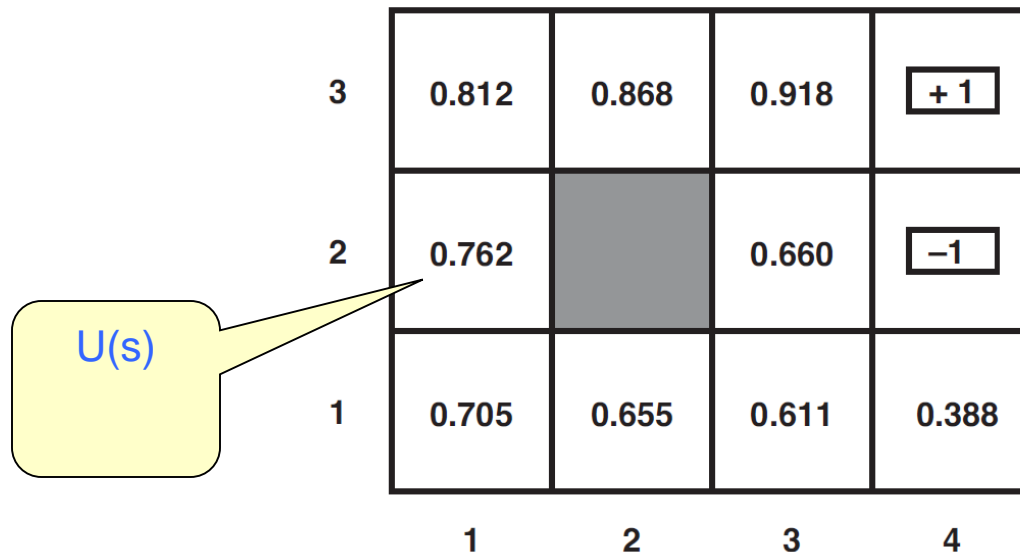
### Technique

Q-learning

Value Learning

# Difference between $R(s)$ and $U(s)$

- $R(s)$  – the “**short-term**” reward for being in state (s)
- $U(s)$  – the “**long-term**” **total** reward from (s) onward




|   |       |       |       |       |
|---|-------|-------|-------|-------|
| 3 | 0.812 | 0.868 | 0.918 | $+1$  |
| 2 | 0.762 |       | 0.660 | $-1$  |
| 1 | 0.705 | 0.655 | 0.611 | 0.388 |
|   | 1     | 2     | 3     | 4     |

$$R(s) = \begin{cases} +1 \\ -0.04 \\ -1 \end{cases}$$

# Given $U(s)$ , compute optimal policy $\pi^*(s)$

- Pick the action with the maximal expected utility (MEU)

$$\pi^*(s) = \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U(s')$$



**Up:**

**Down:**

**Left:**


**Right:**

|   |       |       |       |       |
|---|-------|-------|-------|-------|
| 3 | 0.812 | 0.868 | 0.918 | +1    |
| 2 | 0.762 |       | 0.660 | -1    |
| 1 | 0.705 | 0.655 | 0.611 | 0.388 |
|   | 1     | 2     | 3     | 4     |

# Given $U(s)$ , compute optimal policy $\pi^*(s)$

- Pick the action with the maximal expected utility (MEU)

$$\pi^*(s) = \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U(s')$$



**Up:**

$$0.8 * 0.918 + 0.1 * 0.868 + 0.1 * 1 = 0.9212$$

**Down:**

$$0.8 * 0.660 + 0.1 * 0.868 + 0.1 * 1 = 0.7148$$

**Left:**

$$0.8 * 0.868 + 0.1 * 0.660 + 0.1 * 0.918 = 0.8522$$

**Right:**

$$0.8 * 1 + 0.1 * 0.918 + 0.1 * 0.660 = 0.9578$$

|   |       |       |       |       |
|---|-------|-------|-------|-------|
|   |       |       |       |       |
| 3 | 0.812 | 0.868 | 0.918 | +1    |
| 2 | 0.762 |       | 0.660 | -1    |
| 1 | 0.705 | 0.655 | 0.611 | 0.388 |
|   | 1     | 2     | 3     | 4     |

The problem is that  
"we don't have  $U(s)$ " in  
the first place


# Bellman equation (1957)

---

$$U(s) = R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U(s')$$

Example:

$$U(1,1) = -0.04 + \gamma \max \begin{cases} 0.8U(1,2) + 0.1U(2,1) + 0.1U(1,1), & (Up) \\ 0.9U(1,1) + 0.1U(1,2), & (Left) \\ 0.9U(1,1) + 0.1U(2,1), & (Down) \\ 0.8U(2,1) + 0.1U(1,2) + 0.1U(1,1) \end{cases} \quad (Right)$$



Similarly, write down  $U(1,2)$ ,  $U(1,3)$ ,  $U(1,4)$ ,  
 $U(2,1)$ ,  $U(2,2)$ ,  $U(2,3)$ ,  $U(2,4)$ ,  
 $U(3,1)$ ,  $U(3,2)$ ,  $U(3,3)$ ,  $U(3,4)$

Solutions to these equations are unique!


# Bellman equation (1957)

$$U(s) = R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U(s')$$

You can't directly solve these equations due to the non-linear (**max**) operation

Example:

$$U(1,1) = -0.04 + \gamma \max \begin{cases} 0.8U(1,2) + 0.1U(2,1) + 0.1U(1,1), & (Up) \\ 0.9U(1,1) + 0.1U(1,2), & (Left) \\ 0.9U(1,1) + 0.1U(2,1), & (Down) \\ 0.8U(2,1) + 0.1U(1,2) + 0.1U(1,1) \end{cases} \quad (Right)$$



Similarly, write down  $U(1,2)$ ,  $U(1,3)$ ,  $U(1,4)$ ,

$U(2,1)$ ,  $U(2,2)$ ,  $U(2,3)$ ,  $U(2,4)$ ,

$U(3,1)$ ,  $U(3,2)$ ,  $U(3,3)$ ,  $U(3,4)$

Solutions to these equations are unique!

# Policy evaluation

---

- Given a policy  $\pi(s)$ , compute  $U(s)$

$$U_i(s) = R(s) + \gamma \sum_{s'} P(s' | s, \pi_i(s)) U_i(s')$$

This is a set of *linear* equations, which is polynomial time solvable (by LP solver); in contrast, the Bellman equation  $U(s) = R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U(s')$  is nonlinear.




# Policy evaluation: *example*

- Given a policy  $\pi(s)$ , compute  $U(s)$

But you do not know  
the **transition model**  
 $P(s' | s, a)$

$$U_i(s) = R(s) + \gamma \sum_{s'} P(s' | s, \pi_i(s)) U_i(s')$$

This is a set of *linear* equations, which is polynomial time solvable (by LP solver); in contrast, the Bellman equation  $U(s) = R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U(s')$  is nonlinear.



$$\begin{aligned} U_i(1,1) &= -0.04 + 0.8U_i(1,2) + 0.1U_i(1,1) + 0.1U_i(2,1) , \\ U_i(1,2) &= -0.04 + 0.8U_i(1,3) + 0.2U_i(1,2) , \\ &\vdots \end{aligned}$$

You could have directly  
solved these equations

# Reinforcement learning

---

- Passive RL
  - The policy is fixed, but the transition model is unknown, and we want to learn the utility  $U(s)$
- Active RL
  - Need to learn a policy (i.e., deciding what actions to take)

# Reinforcement learning

---

- Passive RL
  - The policy is fixed, but the transition model is unknown, and we want to learn the utility  $U(s)$
- Active RL
  - Need to learn a policy (i.e., deciding what actions to take)

# Passive Reinforcement Learning

- Objective
  - Assume policy  $\pi(s)$  is fixed, but transition model is unknown, and we would like to directly learn the utility  $U(s)$


|   |   |   |   |    |
|---|---|---|---|----|
| 3 | → | → | → | +1 |
| 2 | ↑ |   | ↑ | -1 |
| 1 | ↑ | ← | ← | ←  |
|   | 1 | 2 | 3 | 4  |

How?

|   |       |       |       |       |
|---|-------|-------|-------|-------|
| 3 | 0.812 | 0.868 | 0.918 | +1    |
| 2 | 0.762 |       | 0.660 | -1    |
| 1 | 0.705 | 0.655 | 0.611 | 0.388 |
|   | 1     | 2     | 3     | 4     |

# Passive Reinforcement Learning

- Objective
  - Assume policy  $\pi(s)$  is fixed, but transition model is unknown, and we would like to directly learn the utility  $U(s)$




- Executes a set of trials

$(1, 1) \xrightarrow{.04} (1, 2) \xrightarrow{.04} (1, 3) \xrightarrow{.04} (1, 2) \xrightarrow{.04} (1, 3) \xrightarrow{.04} (2, 3) \xrightarrow{.04} (3, 3) \xrightarrow{.04} (4, 3) +1$   
 $(1, 1) \xrightarrow{.04} (1, 2) \xrightarrow{.04} (1, 3) \xrightarrow{.04} (2, 3) \xrightarrow{.04} (3, 3) \xrightarrow{.04} (3, 2) \xrightarrow{.04} (3, 3) \xrightarrow{.04} (4, 3) +1$   
 $(1, 1) \xrightarrow{.04} (2, 1) \xrightarrow{.04} (3, 1) \xrightarrow{.04} (3, 2) \xrightarrow{.04} (4, 2) -1$  .

# Direct utility estimation [Widrow & Hoff, 1960]

- Learn a map from states to utilities

- $(1,1) \rightarrow$
- $(1,2) \rightarrow$
- $(1,3) \rightarrow$
- ...




- Executes a set of trials

$(1,1) \xrightarrow{.04} (1,2) \xrightarrow{.04} (1,3) \xrightarrow{.04} (1,2) \xrightarrow{.04} (1,3) \xrightarrow{.04} (2,3) \xrightarrow{.04} (3,3) \xrightarrow{.04} (4,3) +1$   
 $(1,1) \xrightarrow{.04} (1,2) \xrightarrow{.04} (1,3) \xrightarrow{.04} (2,3) \xrightarrow{.04} (3,3) \xrightarrow{.04} (3,2) \xrightarrow{.04} (3,3) \xrightarrow{.04} (4,3) +1$   
 $(1,1) \xrightarrow{.04} (2,1) \xrightarrow{.04} (3,1) \xrightarrow{.04} (3,2) \xrightarrow{.04} (4,2) -1$

# Direct utility estimation [Widrow & Hoff, 1960]

- Learn a map from states to utilities

- $(1,1) \rightarrow 0.76$
- $(1,2) \rightarrow$
- $(1,3) \rightarrow$
- ...




- Executes a set of trials

$(1,1) \xrightarrow{.04} (1,2) \xrightarrow{.04} (1,3) \xrightarrow{.04} (1,2) \xrightarrow{.04} (1,3) \xrightarrow{.04} (2,3) \xrightarrow{.04} (3,3) \xrightarrow{.04} (4,3) +1$   
 $(1,1) \xrightarrow{.04} (1,2) \xrightarrow{.04} (1,3) \xrightarrow{.04} (2,3) \xrightarrow{.04} (3,3) \xrightarrow{.04} (3,2) \xrightarrow{.04} (3,3) \xrightarrow{.04} (4,3) +1$   
 $(1,1) \xrightarrow{.04} (2,1) \xrightarrow{.04} (3,1) \xrightarrow{.04} (3,2) \xrightarrow{.04} (4,2) -1$  .

# Direct utility estimation [Widrow & Hoff, 1960]

- Learn a map from states to utilities
  - $(1,1) \rightarrow 0.72$
  - $(1,2) \rightarrow (0.76 + 0.84) / 2 = 0.80$
  - $(1,3) \rightarrow$
  - ...



- Executes a set of trials


$(1,1) \xrightarrow{.04} (1,2) \xrightarrow{.04} (1,3) \xrightarrow{.04} (1,2) \xrightarrow{.04} (1,3) \xrightarrow{.04} (2,3) \xrightarrow{.04} (3,3) \xrightarrow{.04} (4,3) +1$   
 $(1,1) \xrightarrow{.04} (1,2) \xrightarrow{.04} (1,3) \xrightarrow{.04} (2,3) \xrightarrow{.04} (3,3) \xrightarrow{.04} (3,2) \xrightarrow{.04} (3,3) \xrightarrow{.04} (4,3) +1$   
 $(1,1) \xrightarrow{.04} (2,1) \xrightarrow{.04} (3,1) \xrightarrow{.04} (3,2) \xrightarrow{.04} (4,2) -1$  .



# Direct utility estimation [Widrow & Hoff, 1960]

- Learn a map from states to utilities

- $(1,1) \rightarrow 0.72$
- $(1,2) \rightarrow (0.76 + 0.84) / 2 = 0.80$
- $(1,3) \rightarrow (0.80 + 0.88) / 2 = 0.84$
- ...




- Executes a set of trials

$(1,1) \xrightarrow{.04} (1,2) \xrightarrow{.04} (1,3) \xrightarrow{.04} (1,2) \xrightarrow{.04} (1,3) \xrightarrow{.04} (2,3) \xrightarrow{.04} (3,3) \xrightarrow{.04} (4,3) +1$   
 $(1,1) \xrightarrow{.04} (1,2) \xrightarrow{.04} (1,3) \xrightarrow{.04} (2,3) \xrightarrow{.04} (3,3) \xrightarrow{.04} (3,2) \xrightarrow{.04} (3,3) \xrightarrow{.04} (4,3) +1$   
 $(1,1) \xrightarrow{.04} (2,1) \xrightarrow{.04} (3,1) \xrightarrow{.04} (3,2) \xrightarrow{.04} (4,2) -1 .$

# Direct utility estimation [Widrow & Hoff, 1960]

- Learn a map from states to utilities

- $(1,1) \rightarrow 0.76$
- $(1,2) \rightarrow (0.76 + 0.84) / 2 = 0.80$
- $(1,3) \rightarrow (0.80 + 0.88) / 2 = 0.84$
- ...



**Problem:** misses important information -  
utilities of states are *NOT* independent!


$$U^\pi(s) = R(s) + \gamma \sum_{s'} P(s' | s, \pi(s)) U^\pi(s')$$

- Executes a set of trials

$(1,1) \cdot_{.04} \rightsquigarrow (1,2) \cdot_{.04} \rightsquigarrow (1,3) \cdot_{.04} \rightsquigarrow (1,2) \cdot_{.04} \rightsquigarrow (1,3) \cdot_{.04} \rightsquigarrow (2,3) \cdot_{.04} \rightsquigarrow (3,3) \cdot_{.04} \rightsquigarrow (4,3) +1$   
 $(1,1) \cdot_{.04} \rightsquigarrow (1,2) \cdot_{.04} \rightsquigarrow (1,3) \cdot_{.04} \rightsquigarrow (2,3) \cdot_{.04} \rightsquigarrow (3,3) \cdot_{.04} \rightsquigarrow (3,2) \cdot_{.04} \rightsquigarrow (3,3) \cdot_{.04} \rightsquigarrow (4,3) +1$   
 $(1,1) \cdot_{.04} \rightsquigarrow (2,1) \cdot_{.04} \rightsquigarrow (3,1) \cdot_{.04} \rightsquigarrow (3,2) \cdot_{.04} \rightsquigarrow (4,2) -1 .$

# Direct utility estimation [Widrow & Hoff, 1960]

- Learn a map from states to utilities
  - $(1,1) \rightarrow 0.76$
  - $(1,2) \rightarrow (0.76 + 0.84) / 2 = 0.80$
  - $(1,3) \rightarrow (0.80 + 0.88) / 2 = 0.84$
  - ...



**Problem:** misses important information -  
utilities of states are *NOT* independent!

Could have updated  $U(1,1)$   
based on previous value of  
 $U(1,2)$  but it didn't...

$$U^\pi(s) = R(s) + \gamma \sum_{s'} P(s' | s, \pi(s)) U^\pi(s')$$

- Executes a set of trials

$(1,1) \cdot_{.04} \rightsquigarrow (1,2) \cdot_{.04} \rightsquigarrow (1,3) \cdot_{.04} \rightsquigarrow (1,2) \cdot_{.04} \rightsquigarrow (1,3) \cdot_{.04} \rightsquigarrow (2,3) \cdot_{.04} \rightsquigarrow (3,3) \cdot_{.04} \rightsquigarrow (4,3) +1$   
 $(1,1) \cdot_{.04} \rightsquigarrow (1,2) \cdot_{.04} \rightsquigarrow (1,3) \cdot_{.04} \rightsquigarrow (2,3) \cdot_{.04} \rightsquigarrow (3,3) \cdot_{.04} \rightsquigarrow (3,2) \cdot_{.04} \rightsquigarrow (3,3) \cdot_{.04} \rightsquigarrow (4,3) +1$   
 $(1,1) \cdot_{.04} \rightsquigarrow (2,1) \cdot_{.04} \rightsquigarrow (3,1) \cdot_{.04} \rightsquigarrow (3,2) \cdot_{.04} \rightsquigarrow (4,2) -1$  .

# Question

---

- While learning  $\mathbf{U}(\mathbf{s})$ , how to leverage state dependency (e.g., Bellman equation) to speed up the iterations?
  - Without it, “**Direct Utility Estimation**” method is slow to converge

# Adaptive dynamic programming (ADP)

- Learn the transition model  $\mathbf{P}(\mathbf{s}' \mid \mathbf{s}, \mathbf{a})$  first
  - **Input:** a “state-action” pair  $(\mathbf{s}, \mathbf{a})$
  - **Output:** a new state  $\mathbf{s}'$
  - Treat it as a table of probabilities
    - Count **how often** each action outcome  $\mathbf{s}'$  occurs, given  $(\mathbf{s}, \mathbf{a})$

$(1, 1) \cdot .04 \rightsquigarrow (1, 2) \cdot .04 \rightsquigarrow (1, 3) \cdot .04 \rightsquigarrow (1, 2) \cdot .04 \rightsquigarrow (1, 3) \cdot .04 \rightsquigarrow (2, 3) \cdot .04 \rightsquigarrow (3, 3) \cdot .04 \rightsquigarrow (4, 3)_{+1}$   
 $(1, 1) \cdot .04 \rightsquigarrow (1, 2) \cdot .04 \rightsquigarrow (1, 3) \cdot .04 \rightsquigarrow (2, 3) \cdot .04 \rightsquigarrow (3, 3) \cdot .04 \rightsquigarrow (3, 2) \cdot .04 \rightsquigarrow (3, 3) \cdot .04 \rightsquigarrow (4, 3)_{+1}$   
 $(1, 1) \cdot .04 \rightsquigarrow (2, 1) \cdot .04 \rightsquigarrow (3, 1) \cdot .04 \rightsquigarrow (3, 2) \cdot .04 \rightsquigarrow (4, 2)_{-1} .$

$P((2, 3) \mid (1, 3), \text{Right})$  is estimated to be  $2/3$



$$U^\pi(s) = R(s) + \gamma \sum_{s'} P(s' \mid s, \pi(s)) U^\pi(s')$$

# Question

---

- How do you count?
  - Often times, the results need to be **normalized**
  - In other words, we want frequency (e.g.,  $P(s' | s, a)$  )

# Example: Expected Age

Goal: Compute expected age of this **CS 360** students

Known  $P(A)$

$$E[A] = \sum_a P(a) \cdot a = 0.35 \times 20 + \dots$$

Without  $P(A)$ , instead collect samples  $[a_1, a_2, \dots, a_N]$

Unknown  $P(A)$ : “Model Based”

$$\hat{P}(a) = \frac{\text{num}(a)}{N}$$

$$E[A] \approx \sum_a \hat{P}(a) \cdot a$$

Why does this work? Because eventually you learn the right model.

Unknown  $P(A)$ : “Model Free”

$$E[A] \approx \frac{1}{N} \sum_i a_i$$

Why does this work? Because samples appear with the right frequencies.

# Question

---

- How to compute the “**moving average**”?
  - Don't want to wait for all samples to come in
  - Need to have the average of samples seen so far



# Exponential Moving Average

- Exponential moving average

- The running interpolation update:  $\bar{x}_n = (1 - \alpha) \cdot \bar{x}_{n-1} + \alpha \cdot x_n$
- Makes recent samples more important:

$$\bar{x}_n = \frac{x_n + (1 - \alpha) \cdot x_{n-1} + (1 - \alpha)^2 \cdot x_{n-2} + \dots}{1 + (1 - \alpha) + (1 - \alpha)^2 + \dots}$$

- Forgets about the past (distant past values were wrong anyway)
- Decreasing learning rate (alpha) can give converging averages


# Temporal-difference learning

---

- Idea: learn from every experience
  - Update  $U(s)$  each time we experience a transition  $(s, a, s', r)$
  - Likely outcome  $(s')$  will contribute more often
- Adjust value based on the value of successor state
  - Moving average

Sample of  $U(s)$ : *sample*  $= R(s) + \gamma U^\pi(s')$

Update to  $U(s)$ :  $U^\pi(s) = (1 - \alpha) U^\pi(s) + (\alpha) \textit{sample}$



$$U^\pi(s) \leftarrow U^\pi(s) + \alpha(R(s) + \gamma U^\pi(s') - U^\pi(s))$$

# Temporal-difference learning (*example*)

---

- Suppose that, after the 1st trial, we have
  - $U(1,3) = \mathbf{0.84}$
  - $U(2,3) = 0.92$
- In the 2<sup>nd</sup> trial, since transition  $(1,3) \rightarrow (2,3)$  occurs again
  - $U(1,3) = -0.04 + U(2,3) = -0.04 + 0.92 = \mathbf{0.88}$
  - Assume ( $\gamma=1$ ) for ease of understanding

$$U^\pi(s) \leftarrow U^\pi(s) + \alpha(R(s) + \gamma U^\pi(s') - U^\pi(s))$$

$(1,1)_{-.04} \rightsquigarrow (1,2)_{-.04} \rightsquigarrow (1,3)_{-.04} \rightsquigarrow (1,2)_{-.04} \rightsquigarrow (1,3)_{-.04} \rightsquigarrow (2,3)_{-.04} \rightsquigarrow (3,3)_{-.04} \rightsquigarrow (4,3)_{+1}$   
 $(1,1)_{-.04} \rightsquigarrow (1,2)_{-.04} \rightsquigarrow (1,3)_{-.04} \rightsquigarrow (2,3)_{-.04} \rightsquigarrow (3,3)_{-.04} \rightsquigarrow (3,2)_{-.04} \rightsquigarrow (3,3)_{-.04} \rightsquigarrow (4,3)_{+1}$   
 $(1,1)_{-.04} \rightsquigarrow (2,1)_{-.04} \rightsquigarrow (3,1)_{-.04} \rightsquigarrow (3,2)_{-.04} \rightsquigarrow (4,2)_{-1}$  .

# Temporal-difference learning (*example*)

- Suppose that, after the 1st trial, we have
  - $U(1,3) = \mathbf{0.84}$
  - $U(2,3) = 0.92$
- In the 2<sup>nd</sup> trial, since transition  $(1,3) \rightarrow (2,3)$  occurs again
  - $U(1,3) = -0.04 + U(2,3) = -0.04 + 0.92 = \mathbf{0.88}$
  - Assume ( $\gamma=1$ ) for ease of understanding

But we could do better

$$U^\pi(s) \leftarrow U^\pi(s) + \alpha(R(s) + \gamma U^\pi(s') - U^\pi(s))$$

**( 0.88 - 0.84 )**

$(1,1)_{-.04} \rightsquigarrow (1,2)_{-.04} \rightsquigarrow (1,3)_{-.04} \rightsquigarrow (1,2)_{-.04} \rightsquigarrow (1,3)_{-.04} \rightsquigarrow (2,3)_{-.04} \rightsquigarrow (3,3)_{-.04} \rightsquigarrow (4,3)_{+1}$   
 $(1,1)_{-.04} \rightsquigarrow (1,2)_{-.04} \rightsquigarrow (1,3)_{-.04} \rightsquigarrow (2,3)_{-.04} \rightsquigarrow (3,3)_{-.04} \rightsquigarrow (3,2)_{-.04} \rightsquigarrow (3,3)_{-.04} \rightsquigarrow (4,3)_{+1}$   
 $(1,1)_{-.04} \rightsquigarrow (2,1)_{-.04} \rightsquigarrow (3,1)_{-.04} \rightsquigarrow (3,2)_{-.04} \rightsquigarrow (4,2)_{-1}$

# Reinforcement learning

---

- Passive RL
  - The policy is fixed, but the transition model is unknown, and we want to learn the utility  $U(s)$
- Active RL
  - Need to learn a policy (i.e., deciding what actions to take)

# Trade-off: *exploitation vs. exploration*

---

- When to explore?
  - **Random actions:** explore a fixed amount
  - **Better idea:** try to explore areas whose badness is not (yet) established, but eventually stop exploring
- Exploration function
  - Takes a value estimate **u** and a visit count **n**, and returns an optimistic utility, e.g.

$$f(u, n) = u + k/n$$

# Q-learning

---

- Learn an “action-utility” representation, denoted  $\mathbf{Q(s,a)}$ 
  - $Q(s,a)$  denotes the value of doing action (a) in state (s)

$$U(s) = \max_a Q(s, a)$$

- **Insight:** Agent with  $\mathbf{Q(s,a)}$  no longer needs  $\mathbf{P(s' | s,a)}$  for action selection
  - Called “model-free” method

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a))$$

# Detour: Q-Value Iteration

- Value iteration: find successive (depth-limited) values

- Start with  $U_0(s)$
- Given  $U_k$ , calculate the depth  $k+1$  values for all states:

$$U^\pi(s) \leftarrow U^\pi(s) + \alpha(R(s) + \gamma U^\pi(s') - U^\pi(s))$$

- But Q-values are more useful, so compute them instead

- Start with  $Q_0(s,a)$
- Given  $Q_k$ , calculate the depth  $k+1$  Q-values for all Q-states:

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a))$$



# Q-Learning Properties

---

- Amazing result: Q-learning converges to optimal policy -- even if you're acting sub optimally!
- This is called off-policy learning
- Caveats:
  - You have to explore enough...
  - In the limit, it doesn't matter how you select actions (!)

# The Story So Far: MDPs and RL

## Known MDP: Offline Solution

### Goal

Compute  $U^*, Q^*, \pi^*$

Evaluate a fixed policy  $\pi$

### Technique

Value / policy iteration

Policy evaluation

## Unknown MDP: Model-Based

### Goal

Compute  $U^*, Q^*, \pi^*$

Evaluate a fixed policy  $\pi$

### Technique

VI/PI on approx. MDP

PE on approx. MDP

## Unknown MDP: Model-Free

### Goal

Compute  $U^*, Q^*, \pi^*$

Evaluate a fixed policy  $\pi$

### Technique

Q-learning

Value Learning